

```
In [2]: import tensorflow as tf
        from tensorflow.keras.preprocessing.image import ImageDataGenerator
```

```
2024-03-24 17:04:45.799744: E external/local_xla/xla/stream_executor/cuda/cuda_dnn.c
c:9261] Unable to register cuDNN factory: Attempting to register factory for plugin
cuDNN when one has already been registered
2024-03-24 17:04:45.799862: E external/local_xla/xla/stream_executor/cuda/cuda_fft.c
c:607] Unable to register cuFFT factory: Attempting to register factory for plugin c
uFFT when one has already been registered
2024-03-24 17:04:45.964900: E external/local_xla/xla/stream_executor/cuda/cuda_blas.
cc:1515] Unable to register cuBLAS factory: Attempting to register factory for plugi
n cuBLAS when one has already been registered
```

```
In [3]: train_datagen = ImageDataGenerator(
        rescale=1./255,
    )

    training_set = train_datagen.flow_from_directory(
        '/kaggle/input/brain-tumour-dataset/Dataset/Training',
        target_size=(224, 224),
        batch_size=64,
        class_mode='categorical')
```

Found 5712 images belonging to 4 classes.

```
In [4]: validation_datagen = ImageDataGenerator(rescale=1./255)

    validation_set = validation_datagen.flow_from_directory(
        '/kaggle/input/brain-tumour-dataset/Dataset/validation',
        target_size=(224, 224),
        batch_size=64,
        class_mode='categorical',
        shuffle = False)
```

Found 1239 images belonging to 4 classes.

Building the CNN :

Customising the AlexNET architecture for better result. Here we have used strides=4 in the 3rd,4th and 5th Convolution layers and omitted the last 3rd Pooling Layer (Max) from the architecture, to achieve more efficiency by experiment. (Which is different from the classic AlexNET model)

```
In [5]: # Initialising the CNN
        cnn = tf.keras.models.Sequential()

        # Convolutional Layer
        cnn.add(tf.keras.layers.Conv2D(filters=96, kernel_size=11, strides=4, padding='same'))
        cnn.add(tf.keras.layers.Lambda(tf.nn.local_response_normalization)) # Normalisation

        # Pooling Layer (Max Pooling)
        cnn.add(tf.keras.layers.MaxPooling2D(pool_size=3, strides=2))

        # 2nd Convolutional Layer
        cnn.add(tf.keras.layers.Conv2D(filters=256, kernel_size=5, strides = 4, padding='sa
```

```

cnn.add(tf.keras.layers.Lambda(tf.nn.local_response_normalization)) # Normalisation

# 2nd Pooling Layer (Max Pooling)
cnn.add(tf.keras.layers.MaxPooling2D(pool_size=3, strides=2))

# Convolutional Layer 3
cnn.add(tf.keras.layers.Conv2D(filters = 384, kernel_size = 3, strides=4, padding='

# Convolutional Layer 4
cnn.add(tf.keras.layers.Conv2D(filters = 384, kernel_size = 3, strides=4, padding='

# Convolutional Layer 5
cnn.add(tf.keras.layers.Conv2D(filters = 256, kernel_size = 3, strides=4, padding='

# Flattening Layer
cnn.add(tf.keras.layers.Flatten())

# Full Connection
cnn.add(tf.keras.layers.Dense(units=4096,activation="relu")) # 1st Hidden Layer
cnn.add(tf.keras.layers.Dropout(0.5))

cnn.add(tf.keras.layers.Dense(units=4096,activation="relu")) # 2nd Hidden Layer
cnn.add(tf.keras.layers.Dropout(0.5))

cnn.add(tf.keras.layers.Dense(units=4,activation="softmax")) # Output Layer

# Compiling the CNN :
cnn.compile(optimizer="adam", loss = "categorical_crossentropy", metrics=["accuracy

```

```

/opt/conda/lib/python3.10/site-packages/keras/src/layers/convolutional/base_conv.py:
99: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When
using Sequential models, prefer using an `Input(shape)` object as the first layer in
the model instead.
super().__init__(

```

Training the Model :

```

In [6]: # Training and Evaluating the CNN :
history = cnn.fit(x = training_set, validation_data = validation_set, epochs=25)

```

Epoch 1/25

```

/opt/conda/lib/python3.10/site-packages/keras/src/trainers/data_adapters/py_dataset_
adapter.py:122: UserWarning: Your `PyDataset` class should call `super().__init__(**
kwargs)` in its constructor. `**kwargs` can include `workers`, `use_multiprocessing
`, `max_queue_size`. Do not pass these arguments to `fit()`, as they will be ignore
d.

```

```

self._warn_if_super_not_called()

```

2/90 ————— 6s 70ms/step - accuracy: 0.3477 - loss: 1.3849

```

WARNING: All log messages before absl::InitializeLog() is called are written to STDERR

```

```


I0000 00:00:1711299928.898236      102 device_compiler.h:186] Compiled cluster using
XLA! This line is logged at most once for the lifetime of the process.

```

```

W0000 00:00:1711299928.920726      102 graph_launch.cc:671] Fallback to op-by-op mode
because memset node breaks graph update


```


23/90  **1:00** 908ms/step - accuracy: 0.2929 - loss: 1.4451


W0000 00:00:1711299948.780342 103 graph_launch.cc:671] Fallback to op-by-op mode
because memset node breaks graph update


89/90  **0s** 674ms/step - accuracy: 0.3138 - loss: 1.3627


W0000 00:00:1711299989.933804 103 graph_launch.cc:671] Fallback to op-by-op mode
because memset node breaks graph update


90/90  **92s** 821ms/step - accuracy: 0.3158 - loss: 1.3591 - val_accuracy: 0.6602 - val_loss: 0.8481
Epoch 2/25


90/90  **24s** 247ms/step - accuracy: 0.6900 - loss: 0.7167 - val_accuracy: 0.7659 - val_loss: 0.7633
Epoch 3/25


90/90  **23s** 237ms/step - accuracy: 0.8262 - loss: 0.5204 - val_accuracy: 0.7869 - val_loss: 0.5701
Epoch 4/25


90/90  **23s** 242ms/step - accuracy: 0.8709 - loss: 0.4036 - val_accuracy: 0.8176 - val_loss: 0.5554
Epoch 5/25


90/90  **24s** 244ms/step - accuracy: 0.8885 - loss: 0.3316 - val_accuracy: 0.8297 - val_loss: 0.4792
Epoch 6/25


90/90  **23s** 240ms/step - accuracy: 0.9060 - loss: 0.2802 - val_accuracy: 0.8386 - val_loss: 0.4303
Epoch 7/25


90/90  **24s** 245ms/step - accuracy: 0.9071 - loss: 0.2718 - val_accuracy: 0.8370 - val_loss: 0.4134
Epoch 8/25


90/90  **23s** 235ms/step - accuracy: 0.9271 - loss: 0.2249 - val_accuracy: 0.8378 - val_loss: 0.3432
Epoch 9/25


90/90  **24s** 252ms/step - accuracy: 0.9350 - loss: 0.1911 - val_accuracy: 0.8797 - val_loss: 0.2910
Epoch 10/25


90/90  **23s** 242ms/step - accuracy: 0.9353 - loss: 0.1712 - val_accuracy: 0.9104 - val_loss: 0.2616
Epoch 11/25


90/90  **23s** 240ms/step - accuracy: 0.9297 - loss: 0.1873 - val_accuracy: 0.8878 - val_loss: 0.3010
Epoch 12/25


90/90  **23s** 240ms/step - accuracy: 0.9470 - loss: 0.1501 - val_accuracy: 0.8878 - val_loss: 0.2847
Epoch 13/25


90/90  **23s** 239ms/step - accuracy: 0.9520 - loss: 0.1190 - val_accuracy: 0.9177 - val_loss: 0.2099
Epoch 14/25


90/90  **23s** 238ms/step - accuracy: 0.9722 - loss: 0.0946 - val_accuracy: 0.9040 - val_loss: 0.2663
Epoch 15/25

90/90  **23s** 243ms/step - accuracy: 0.9649 - loss: 0.1024 - val_accuracy: 0.9266 - val_loss: 0.1841
Epoch 16/25

90/90  **23s** 237ms/step - accuracy: 0.9673 - loss: 0.0950 - val_accuracy: 0.9266 - val_loss: 0.2012
Epoch 17/25

90/90  **23s** 245ms/step - accuracy: 0.9761 - loss: 0.0718 - val_accuracy: 0.8660 - val_loss: 0.3683
Epoch 18/25

90/90  **23s** 239ms/step - accuracy: 0.9579 - loss: 0.1228 - val_accuracy: 0.9330 - val_loss: 0.1827
Epoch 19/25

90/90  **23s** 236ms/step - accuracy: 0.9688 - loss: 0.0862 - val_accuracy: 0.9209 - val_loss: 0.2530

```

Epoch 20/25
90/90 ————— 23s 240ms/step - accuracy: 0.9789 - loss: 0.0690 - val_ac
curacy: 0.9233 - val_loss: 0.2558
Epoch 21/25
90/90 ————— 23s 240ms/step - accuracy: 0.9822 - loss: 0.0530 - val_ac
curacy: 0.9403 - val_loss: 0.1814
Epoch 22/25
90/90 ————— 23s 239ms/step - accuracy: 0.9887 - loss: 0.0402 - val_ac
curacy: 0.9379 - val_loss: 0.2052
Epoch 23/25
90/90 ————— 23s 241ms/step - accuracy: 0.9892 - loss: 0.0368 - val_ac
curacy: 0.9516 - val_loss: 0.2174
Epoch 24/25
90/90 ————— 23s 239ms/step - accuracy: 0.9735 - loss: 0.0889 - val_ac
curacy: 0.9112 - val_loss: 0.2701
Epoch 25/25
90/90 ————— 23s 245ms/step - accuracy: 0.9723 - loss: 0.0830 - val_ac
curacy: 0.9427 - val_loss: 0.2590

```

```

In [21]: test_datagen = ImageDataGenerator(rescale=1./255)

test_set = test_datagen.flow_from_directory(
    '/kaggle/input/brain-tumour-dataset/Dataset/test',
    target_size=(224, 224),
    batch_size=64,
    class_mode='categorical',
    shuffle=False)

```

Found 72 images belonging to 4 classes.

```

In [22]: import numpy as np

y_pred = cnn.predict(test_set)
y_pred = np.argmax(y_pred, axis=1)
y_true = test_set.classes
print(np.concatenate((y_true.reshape(len(y_true),1),y_pred.reshape(len(y_pred),1)),

```

```

/opt/conda/lib/python3.10/site-packages/keras/src/trainers/data_adapters/py_dataset_adapter.py:122: UserWarning: Your `PyDataset` class should call `super().__init__(**kwargs)` in its constructor. `**kwargs` can include `workers`, `use_multiprocessing`, `max_queue_size`. Do not pass these arguments to `fit()`, as they will be ignored.
  self._warn_if_super_not_called()

```



```
[3 3]
[3 3]
[3 3]
[3 3]
[3 3]
[3 3]
[3 3]
[3 3]
[3 3]
[3 3]
[3 3]
[3 3]
[3 3]
[3 3]
[3 3]
[3 3]
[3 3]
[3 1]]
```

Evaluating the Model :

```
In [23]: # Confusion Matrix :
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_true,y_pred)
print(cm)
```

```
[[18  0  0  0]
 [ 0 17  1  0]
 [ 0  0 18  0]
 [ 0  1  0 17]]
```


```
In [24]: # Accuracy :
from sklearn.metrics import accuracy_score
accuracy = accuracy_score(y_true,y_pred)
print(f"The Accuracy of the model is : {accuracy*100:.2f}%")
```

The Accuracy of the model is : 97.22%


Visualisation of the Model, for perfect Evaluation :

```
In [25]: import matplotlib.pyplot as plt

plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.title('Training and Validation Accuracy in AlexNET')
plt.show()
```



```
In [26]: plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.title('Training and Validation Loss in AlexNET')
plt.show()
```

```
In [28]: from tensorflow.keras.models import load_model  
cnn.save("/kaggle/working/Brain_tumour_prediction_alexnet.hdf5")
```

```
In [ ]:
```